

Build Instructions

To build the code run cmd 'make'

Note:

- makefile uses the mpicc compiler.
- makefile uses '-g -Wall' compiler flag and '-fopenmp' linker flag while compiling the file.
- makefile uses absolute path '/fs/project/PAS1653/transform.o' to link the obj file.
- The generated output file is called 'lab4_hierarchical'.

Execution Instructions

- Reserve 5 nodes 28 ppn using qsub command on owens.
- Execute: `"mpirun -np 5 ./lab4_hierarchical <path_to_input_file>"`

Output

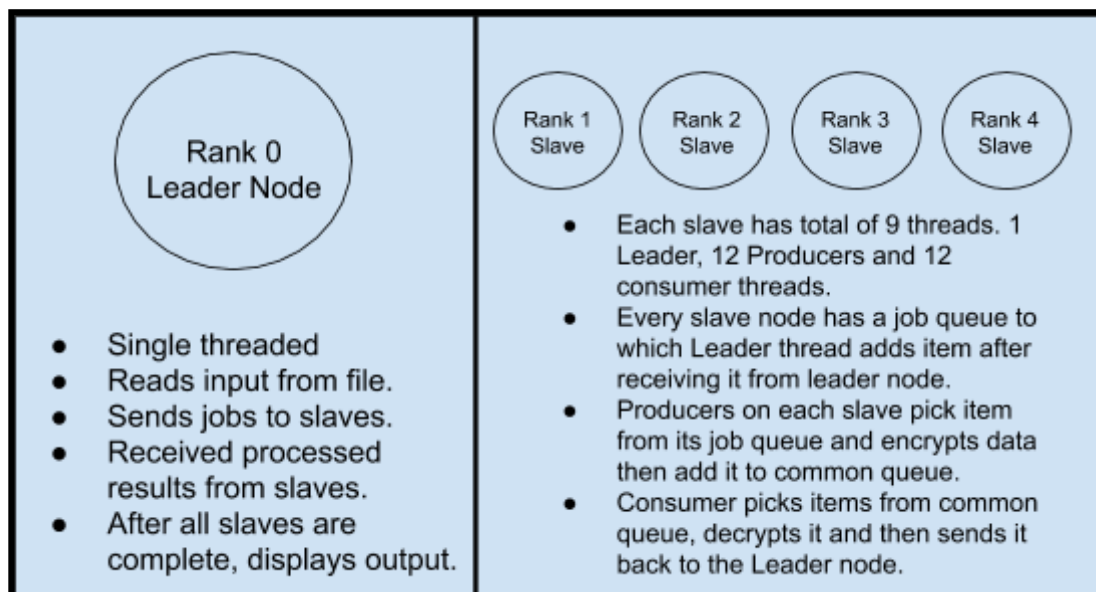
Output format: <position> <command> <p rank> <p tid> <encoded key> <c rank> <c tid> <decoded key>

<position> : the line number not counting invalid input lines

Implementation

Program uses 5 processes to process the input file. Rank 0 process acts as the leader process and is responsible for reading the input file, sending jobs to other processes, collecting the results back and finally displaying the results on the console. Each of the slave processes is running 12 producers and 12 consumer threads that work on the job sent by the leader.

After reading an item from the input file, the leader process adds to the input queue of one of the slave processes. The producer thread on the slave process picks an item from the input queue and calls the transform function to get the encrypted value. After this, the producer adds the item to the common queue from where it is picked by the consumer thread and decrypted. After the consumer thread has completed decryption it sends the final data back to the leader process where it gets added to the output queue. After all slave processes have completed, the leader process serially displays the output on STDOUT and then terminates execution.



Running Times

NOTE: The below runtimes are without the MPI changes that were recommended by Dr. Jones. I tried with the changes and I was seeing significant performance improvement (12P12C/node took 1.54mins instead of 12mins for t1000 test case).

	Lab 4 (MPI)				Lab 3 (OpenMP)			
	Time(2)		Time(1)		Time(2)		Time(1)	
	P	C	Real	User	P	C	Real	User
PC_data_x1	9	38	00: 04.891	00:00.011	2	16	00:01.713	00:13.766
PC_data_t00100	94	141	00:11.323	00:00.017	12	21	00:02.462	00:40.880
PC_data_t01000	910	964	01:03.733	00:00.017	148	162	00:10.662	04:24.239
PC_data_t05000	4699	4756	05:05.679	00:00.017	749	763	00:48.497	21:55.943
PC_data_t10000	9336	9393	09:52:996	00:00.012	1498	1514	01:35.085	43:32.081

NOTE: Producer/Consumer time is in seconds and Real/User time is in mm:ss.ms format.

	Lab 4 (MPI)				Lab 2 (PThreads)			
	Time(2)		Time(1)		Time(2)		Time(1)	
	P	C	Real	User	P	C	Real	User
PC_data_x1	9	38	00: 04.891	00:00.011	3	32	00:02	0:4.192
PC_data_t00100	94	141	00:11.323	00:00.017	15	32	00:02	0:27.158
PC_data_t01000	910	964	01:03.733	00:00.017	143	160	00:11	4:1.991
PC_data_t05000	4699	4756	05:05.679	00:00.017	736	750	00:47	20:27.571
PC_data_t10000	9336	9393	09:52:996	00:00.012	1456	1470	01:32	40:30.495

NOTE: Producer/Consumer time is in seconds and Real/User time is in mm:ss.ms format.

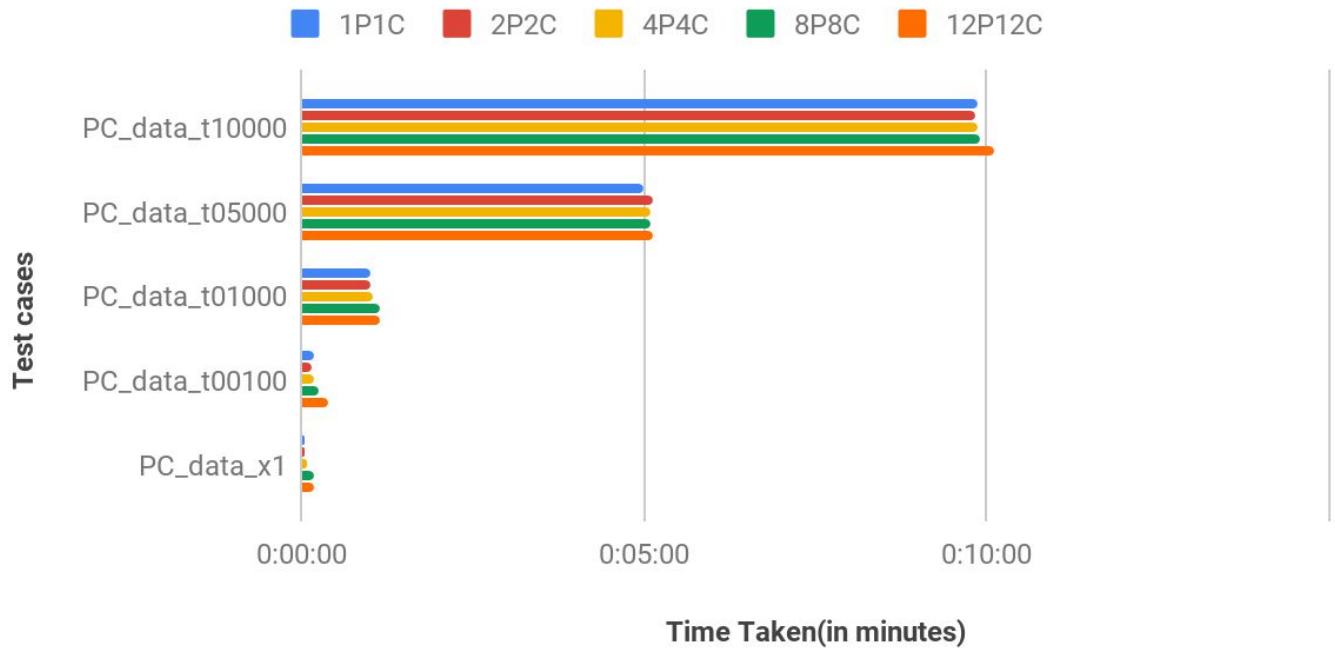
	Lab 4 (MPI)				Lab 1 (Serial)			
	Time(2)		Time(1)		Time(2)		Time(1)	
	P	C	Real	User	P	C	Real	User
PC_data_x1	9	38	00: 04.891	00:00.011	2	1	0:03.585	0:03.513
PC_data_t00100	94	141	00:11.323	00:00.017	10	12	0:22.411	0:22.341
PC_data_t01000	910	964	01:03.733	00:00.017	110	95	3:26.070s	3:25.995
PC_data_t05000	4699	4756	05:05.679	00:00.017	518	535	17:33.385	17:33.147
PC_data_t10000	9336	9393	09:52:996	00:00.012	1045	1043	34:48.672	34:48.429

NOTE: Producer/Consumer time is in seconds and Real/User time is in mm:ss.ms format.

Scalability

Lab 4: Producer Consumer Multithreaded (using MPI)

Time(1): "Real" User Time





Graph1: Performance of MPI implementation for various thread (per slave node) counts.

Lab 4: Producer Consumer Multithreaded (using MPI)

"Real" User Time (vs serial) [faster by]



Graph2: Comparison of MPI implementation (different thread counts) with serial implementation.**Graph3:** Comparison of MPI (different thread counts) with pThread(16P16C) implementation.**Graph3:** Comparison of MPI (different thread counts) with OpenMP(16P16C) implementation.

	Serial	pThread	OpenMP	1P1C	2P2C	4P4C	8P8C	12P12C
PC_data_t10000	34:48.672	01:32.050	01:35.085	09:53.082	09:50.826	09:52.996	09:55.026	10:06.447
PC_data_t05000	17:33.385	00:47.439	00:48.497	05:00.274	05:06.836	05:05.679	05:05.550	05:07.477
PC_data_t01000	03:26.070	00:10.632	00:10.662	01:01.095	01:00.669	01:03.733	01:08.810	01:09.874
PC_data_t00100	00:22.411	00:02.237	00:02.462	00:10.990	00:09.623	00:11.323	00:15.406	00:22.803
PC_data_x1	00:03.585	00:01.910	00:01.713	00:03.885	00:03.691	00:04.891	00:11.903	00:12.078

NOTE: The above data is of Real time returned by Time(1) and is in mm:ss.ms format.

Unexpected Results

- For the same level of threading [(4 Producer, 4 Consumer)X4 Nodes], the time taken by the program to complete is much slower when compared to pThread(16 Producer, 16 Consumer) or OpenMP (16 Producer, 16 Consumer).